

// nasscom.research

The research, in one page.

What we presented at NASSCOM: making
AI evaluation of subjective answers
reproducible, auditable, and accountable.
Swipe through.

// the.problem

01

There is no answer key.

For subjective answers, two graders can disagree — and "the AI agrees with the teacher" only proves resemblance, not that either is right. Neither the teacher nor the machine is the truth.

```
// invariant.evaluator
```

02

<1%

of a score's variance comes from which AI model marks it. Swap the model, the read barely moves. That is reproducibility — not correctness. An invariant reading can still be wrong, which is why a human stays final.

// deterministic.referee

03

A fixed rule, not a model's whim.

A rule-based cascade judges any mark against the scheme and returns the same evidence-linked verdict every run — a stated verdict on 82% of cases, the hardest 18% referred to a human. Nothing changes a mark without a signature.

// the.blind.spot

04

It catches the error no one else can.

When a teacher and an AI agree, nobody looks again — but agreement can be jointly wrong. Reading value-point by value-point, the panel catches marks both graders missed.

// no.self.preference

05

It doesn't take its own side.

The natural fear: an AI protecting its own kind. So we dropped the AI's own model family from the panel — and the verdicts didn't change. It judges the answer, not the machine.

// honest.scope

06

What we can, and cannot, claim.

Demonstrated mechanisms on a small, English-medium corpus — not national rates. Conformance to a marking scheme, not accuracy against absolute truth. A human-expert check is the defined next step.

// the.layer

Reproducible. Auditable. Human-final.

The evaluation layer competency-based assessment was always going to need.
This is what we are building.