

// ai.as.adjudicator

The research, in findings.

What we tested, and what we found — the work we presented at NASSCOM. Swipe through.

// the.corpus

setup

56 answer-cases.
9 Class XII books.

Real handwritten answers, marked against a scheme derived from the board's own marking policy. Not a benchmark — actual student scripts.

```
// finding.01 - invariance
```

<1%

of a score's variance came from which AI model marked it. Swap the model, the read barely moves. Reproducibility — not a claim of correctness.

// finding.02 – the referee

82% / 18%

A deterministic referee returned a stated verdict on 82% of cases and sent the hardest 18% to a human, evidence attached. Same verdict every run.

// finding.03 — no self-preference

03

It doesn't take its own side.

Drop our own model family out of the panel and the verdicts don't change. Four model families, four labs. It judges the answer, not the machine that wrote it.

// honest.scope

What we claim — and what we don't.

Demonstrated mechanisms on a small, English-medium corpus. Conformance to a marking scheme, not accuracy against absolute truth. A human stays the final authority. Reproducible. Auditable. Human-final.