

// ai.as.adjudicator

The research, in findings.

What we tested, and what we found — the work we presented at NASSCOM. Swipe through.

// the.corpus


setup

56 answer-cases.
9 Class XII books.

Real handwritten answers, marked against a scheme derived from the board's own marking policy. Not a benchmark — actual student scripts.

// where.the.mark.comes.from

Change the AI. The mark barely moves.



which AI model marks it — <1% ↑


- 80% — the student's answer itself
- ~19% — other factors (question, interactions)
- <1% — which of 4 AI models marks it

Swap the model across four families and the reading barely moves. That's reproducibility — not accuracy.

// the.deterministic.cascade

82% resolved by rule. 18% to a human.

56 answer-cases adjudicated (62 credited - 6 inadmissible)



Who the rule found for, in the 46 resolved:



The same verdict every run, evidence attached.
Integrity guardrail fired 0x.

CrazyGoldFish of 56 cases • a mechanism, not a population rate

```
// finding.03 — no self-preference
```

03

It doesn't take its own side.

Drop our own model family out of the panel and the verdicts don't change. Four model families, four labs. It judges the answer, not the machine that wrote it.

// honest.scope

What we claim — and what we don't.

Demonstrated mechanisms on a small, English-medium corpus. Conformance to a marking scheme, not accuracy against absolute truth. A human stays the final authority. Reproducible. Auditable. Human-final.