

// who.checks.the.evaluator

# 5 principles for an AI you can hold to account.

When an AI grades the exam, who checks  
the AI?

swipe →

// the.problem

# Agreeing with the teacher is not the same as being right.

If the teacher is wrong and the AI matches, the student still loses. **A wrong mark, twice.** So neither the teacher nor the AI can be the reference.

# 1

## **Judge against the policy, not a person.**

The scheme is derived from the board's own marking policy. The same rule for every student, **not one grader's taste.**

# 2

## Only invariant evaluators get a seat.

An evaluator joins the panel only if its marking barely depends on which model runs it. **A reproducible reference, not a single opinion.**

# 3

## **Decide by a fixed rule, not a model's whim.**

A deterministic four-step cascade runs on the panel's votes. Same votes in, same verdict out, every time. **No model decides the score at the end.**

# 4

## **No evidence, no credit.**

Every mark cites the exact line in the student's answer it rests on. **Auditable, not a black box.**

# 5

## **Judgment stays human.**

When the panel genuinely cannot decide, it escalates to a human with the evidence attached. **Nothing changes a student's mark without a human signing it.**

// the.takeaway

# From grading by opinion to grading you can hold to account.

This is the layer we are building for  
education.