

// what.surprised.us

3 things that surprised me building an AI adjudicator.

We set out to check AI graders. The results changed how I think about it. Swipe.

// surprise.01

01

It didn't defend its own kind.

I expected an AI panel to quietly favour answers its own model family liked. We dropped our family out entirely — the verdicts didn't move. No self-preference.

// the.shared.blind.spot

Both said 4 out of 4. Two entries were never written.

The 8 required entries – 6 attempted, 2 absent:



Teacher	4 / 4
Our AI	4 / 4
Value-point panel	3 / 4

Reading value-point by value-point, the panel caught the two missing entries neither grader had docked.

illustrative case • Accountancy Q3 • expected credit 2.69 → a defensible 3 • mechanism, not a rate
CrazyGoldFish

// surprise.03

03

The model barely mattered.

Change which AI marks the answer, and less than 1% of the score's variance comes from that choice. The answer drives the score, not the model. That is what an invariant reading means.

// the.takeaway

None of it is "the AI is right."

It's that the reading is reproducible, the verdict is auditable, and a human still signs off. That is not the end of the argument. It's a foundation you can finally build on.