

// the.journey

How we got to one question: who checks the evaluator?

Three years, and a few wrong turns that taught us more than the right ones. Swipe.

// step.01

01

We started by building an AI grader.

The pitch was obvious — save teachers time. It worked. Then a teacher asked the question that broke it open.

// step.02

02

"How do you know your AI is right?"

We had been checking whether the AI agreed with the teacher. But if the teacher can be wrong — and they can — agreement proves nothing. There was no ground truth.

// step.03

03

So we tried a human adjudicator.

A person judging disputed marks against the scheme. It worked — but it did not scale, it was still one tired human, and it only ever looked at the disagreements.

// step.04 — where it led

Which left the scariest cases **unopened**.

When a teacher and an AI agree, nobody looks again — even when both are wrong. To catch that, the referee had to read every case. And to do that at scale, it had to become reproducible and auditable. That is the layer we are building.